

A PRACTICAL GUIDE TO AI AGENTS

Building AI Agents from Scratch

A Hands-On Guide with Python



Yanfei Li

Building AI Agents from Scratch

A Hands-On Guide with Python

Yanfei Li

First edition · 2026

Building AI Agents from Scratch

A Hands-On Guide with Python

Copyright © 2026 Yanfei Li. All rights reserved.

First edition, 2026.

The book text and cover may not be reproduced or redistributed without permission, except as permitted by applicable law. Companion source code is distributed under the MIT License included in the code package.

Product names and trademarks belong to their respective owners. This book is independently published and is not endorsed by the providers discussed.

Examples are educational. Some are study fragments; others are complete programs with documented dependencies. Live integrations require your own accounts and may incur charges. Model identifiers and APIs can change.

Code and technical references were reviewed in September 2026. Local checks do not verify every external integration or production deployment.

Companion code

The companion package follows the chapter folders named throughout this book. Refer to the download instructions provided with your purchased edition. Start with the package README and Chapter 3. Code filenames shown in the chapters refer to that package.

About the Author

Yanfei Li, PhD

Yanfei Li is a Senior AI Engineer with a background in energy systems and software development. His work focuses on building AI systems and applying computational methods to solve engineering problems.

He earned a Master of Science in Computer Science with a focus on artificial intelligence from the Georgia Institute of Technology in 2026. He also holds a PhD in Mechanical Engineering, with a focus on modeling, control, and optimization.

Through practical, everyday examples, he aims to make AI agent development accessible and help readers turn ideas into working software.

1 What Are AI Agents?

Suppose a user asks, “What is the hotel limit in our travel policy, and are exceptions allowed?” An agent can search the available documents, read the relevant policy and answer from what it finds. Building that application means connecting a decision process to real operations and carrying their results into the next decision.

What you will learn. Distinguish agents, workflows and chatbots; identify the components of an agent system; follow a request end to end; run a small demonstration of the execution loop.

What you need. Basic Python: functions, dictionaries, loops and running a file in a terminal. Chapter 3 covers environment setup. The demonstration uses local data and needs no model account. Companion folder: `code-examples/chapter-01`.

1.1 Choose the Right Kind of Application

An **agent** observes an environment and chooses actions toward a goal. The environment might be files, business services or physical devices. Russell and Norvig describe this perception–action relationship and tie action selection to a performance measure [1]. For a developer, the useful starting questions are: what can the system observe, what can it do, and what counts as success?

This book builds **LLM agents**: applications in which a large language model selects some next actions from the task and previous results. An **LLM** is a trained model that processes and generates language. The surrounding application supplies context, exposes tools and executes permitted requests.

Table 1.1: Agent, chatbot, script and workflow distinctions

Term	What the term describes	Useful when
Automation script	Code-defined operations and rules.	The required procedure is known and testable.
Workflow	The steps and branches of a process.	The process can be specified explicitly.
LLM workflow	A workflow with model calls at defined points.	A fixed process needs language interpretation or generation.
LLM agent	A system with some model-selected next actions.	Follow-up work depends on varied observations.
Chatbot	A conversational user interface.	People need to ask, clarify, or instruct in conversation.

These terms overlap. A chatbot can expose an agent through a conversation; the same agent can run from a command line or an API. A Python script can

implement either fixed rules or a model-driven loop. Choose the interface and execution design separately.




Figure 1.1: Interface and execution are separate choices. A chat interface can sit over either a fixed workflow or a model-directed loop; a command-line interface can do the same.

For the policy question, compare three implementations. A search script returns matching lines. An LLM workflow follows a fixed search-and-summarize sequence. An agent can use a search result to choose which document to read next.




Figure 1.2: Three execution designs for the same document question. In the bottom design, returned evidence can trigger another model-selected operation. The other designs can still contain explicitly programmed error handling.

Anthropic’s engineering guide recommends starting with simple workflows and adding agent control when a task needs flexible follow-up [2]. This gives us a practical selection rule: use known rules for predictable work, a fixed model step for language processing, and an agent loop when observations should influence the next action.

1.2 A General Architecture for Agent Systems

An **architecture** describes components and their connections. The following reference architecture supports one agent or a team. Several components can live in one Python process.




Figure 1.3: A connected agent system. The top row follows a request from task intake to a checked outcome. Agent 1 is expanded to show how context sources, model access and tools connect to its execution loop; other agents use the same structure. Two arrowheads indicate requests and returns. Numbers point to the implementation chapters.

Read the top row first: the application assigns work to one agent or a team, collects results and returns an outcome. Then follow Agent 1 downward. Instructions, state and selected evidence feed its context. The runtime asks the model for a proposal, checks permissions and limits, and runs a tool if needed. It records the result and checks completion. The green return arrow adds new observations to the next iteration; a finished result goes back to task coordination. Run events support monitoring and evaluation.

A **multi-agent system** coordinates multiple agent instances. Each instance has its own instructions, working context and available actions; instances may share a model service. An **orchestrator** routes work and manages dependencies or handoffs. Use a team when separate capabilities, parallel work or independent review offer a measurable benefit. Chapters 10–13 develop those choices.

The component map below explains each responsibility, why it is useful and where to implement it. Start with the task, model, runtime and tools in Chapters 2–5; add the other capabilities when the application needs them.

Table 1.2: System modules, their purpose and chapter map

Module and meaning	Why or when it is needed	Ch.
Model: a service that proposes text or operations from context	Language interpretation and flexible next-step selection; choose by task tests, limits and cost	2
Development environment: Python, dependencies and configuration	Makes the same program installable and its failures diagnosable	3
Intake, runtime and outcome: accept a goal, execute the loop, report how it ended	Every application needs a defined request, execution boundary and stopping condition	4
Tool contract: a named operation with validated inputs and results	Connect decisions to real reads, calculations or changes; grant only needed operations	5
MCP adapter and skill: a shared tool protocol and reusable procedural instructions	Use MCP to connect compatible applications/services; use a skill to teach a recurring procedure	6
State, context and saved memory: retained run data, selected model input and reusable records	Track progress; add durable storage when work must survive a restart or be reused	7
Reasoning and planning: choose actions and arrange dependencies	Use observations to adapt; add explicit plans when unfinished dependencies become hard to track	8
Retrieval and index: select evidence from prepared sources	Ground answers in material too large or changeable to include in every request	9
Team and orchestrator: distinct agents plus assignment and termination rules	Test when separate roles, permissions or parallel work have a measurable benefit	10

Module and meaning	Why or when it is needed	Ch.
Message contract: a defined envelope for task and result exchange	Prevent confusion about sender, task, correlation, retries and completion	11
Framework: reusable runtime or coordination machinery	Adopt when its tested behavior reduces work for your required control flow	12
Delegation and human review: assign bounded responsibility or ask an authorized person	Handle larger tasks or decisions that need human judgment or permission	13
Evaluation and traces: task tests and records of what ran	Detect unsupported answers and regressions; a reviewer agent alone cannot prove correctness	14
Execution controls: permissions, validation, budgets and failure handling	Apply at actual execution boundaries; decide what may run and what happens when it fails	15
Deployment and operations: API, jobs, workers, storage and release procedure	Make work available beyond one terminal and recover interrupted or failed runs	16
Streaming: incremental delivery of text and progress events	Help an interface show progress; partial output must not trigger an incomplete tool call	17
Training: an offline experiment that changes model parameters	Consider only after diagnosing a learnable model behavior and defining held-out evaluation	18
Multimodal intake: text, images or document pages as evidence	Add when required facts live in more than text; check coverage and extraction separately	19
Database access: scoped queries against structured records	Use original typed records when the task needs filters, joins or aggregates	20
Voice interface: speech input/output around the agent loop	Add when speaking helps the task; test recognition, audio devices and interruption separately	21

The diagram's shared services are reusable capabilities. For example, a document agent may need retrieval, while a calculator agent may need only one tool. Model access, stored memory, retrieval and tool execution remain separate responsibilities even when a framework packages them together.

1.2.1 Inside One Agent




Figure 1.4: General architecture of an LLM agent. Application code prepares inputs and context, calls the model, controls tool execution, records observations, and checks outcomes. Optional saved memory supports later runs. Files, services, and devices remain in the external environment.

The **runtime** is the application code that manages the loop: prepare model input, inspect the proposed action, execute a tool, record its result and decide whether to continue. A **tool** is a named operation the runtime can execute, such as reading a file or querying an API. The model receives a description of that operation and proposes arguments.

State is the information retained during the task: progress, observations and remaining limits. **Context** is the selected information sent in one model request. Saved **memory** makes information available for later work. These distinctions matter because a large stored history need not be sent to the model on every step.

A **goal** states the desired outcome. Instructions and constraints tell the agent how to work toward it. A plan can track dependent steps, but a separate planning component is optional. Anthropic’s guide describes tools, retrieval and memory as capabilities around the model; its recommendation is to choose them for the task [2].

Feedback is information returned after an action, such as retrieved text or an error. **Verification** compares a result with the task’s requirements. For the policy assistant, check the amount, unit, tax condition, exception rule and